

Data Analysis and Knowledge Discovery

Model validation

Jukka Heikkonen

University of Turku
Department of Information Technology

Jukka.Heikkonen@utu.fi


Where are we now

- Using a machine learning method (such as k-nearest neighbours) we can learn a classifier or regressor
- some classification performance measures
 - misclassification rate
 - misclassification rate with costs
 - area under ROC curve (AUC)
- some regression performance measures
 - mean squared error
 - mean absolute error
 - correlation measures also applicable (see early lectures)
- One way or another, all these measures compute how well predicted outputs match true outputs. But what data should we use to compute these?


Where are we now

- Evaluation: given a model, estimate how well it would predict on new data (arriving from similar source as the training data)
- True error rate: predictors expected error rate on the whole population
- Remember: fit to training data does not guarantee generalization to new data
- a complex model may overfit
 - low error on training set, but captures poorly (or not at all) a real predictive pattern
 - when tested on new data, overfitted model predicts inaccurately (or even randomly)
- analogy to human learning: training data is $2 + 2 = 4$, $3 + 2 = 5$ and $4 + 11 = 15$
 - overfitting: the model predicts these correctly, but has no idea what $1 + 1$ equals, since it was not in training set
 - learning: a good model captures the underlying pattern, and thus knows that $1 + 1 = 2$


True model



Reasonable model (15-nearest neighbour)



Overfitting (1-nearest neighbour)



How to choose the best possible k ?

- 1-nearest neighbour classifier/regressor has always 0 training error
- model memorizes all training data, and thus knows the correct output to predict for these instances
- this guarantees nothing about generalization to new data
- a general phenomenon in machine learning, also linear models, neural networks etc. may overfit
- thus models should be tested on **independent test data**

- Validation means measuring the predictor's behavior on data points other than those in the training set
- We need this in order to reliably select the hyperparameters of our model (such as number of neighbours), or in order to choose between models produced by different algorithms
- Also needed for evaluating how well the final model can be expected to generalize to completely new data

Training and test data

To avoid overfitting and to have a more realistic estimation for the error that the model will make for new data, the data set is often split into **training data** and **test data**.


- The model is computed (“trained”) only with the training data.
- The test data set is used to evaluate the model (for instance, to calculate the misclassification rate).
- Sometimes known as the **holdout method**

Typically, $2/3$ of the data are used for training and $1/3$ for testing.

Instances are assigned randomly to the training and the test data.

For classification problems, **stratification** is often applied, guaranteeing that the training and test data have the distribution of classes as the original data.

The holdout method



- In practice, once the model performance has been evaluated, combine all data and use it to train the final model in order to construct as accurate predictor as the data allows.
- Pessimistic bias, the model evaluated, that was trained only on the training set, can be less accurate than the final model trained on all data

The holdout method

- Using a separate test set gives an unbiased estimate of the true expected error (or AUC, or...) on new test data
- Can still be unreliable on too small test sets due to variance, we might get very "lucky" or "unlucky" in which instances happen to fall in the test set
 - you may compute confidence intervals using standard statistical approaches, not covered in this course
- two conflicting goals
 - we want as much data as possible for training in order to be able to discover the possible pattern as well as possible
 - we want as much data as possible for testing in order to get as reliable estimate of true error rate as possible
- when working with Big Data not a problem, we have as much training and test data as we want

Limitations of the holdout method

- What if we have only a data set of 100 instances
- Using e.g. 70 for training and 30 for testing, we get a very unreliable estimate of test performance (too small test set)
- Using e.g. 30 for training and 70 for testing does not help much, now too little training data
- especially bad on imbalanced data, where we could have e.g. 90 instances from majority and only 10 from minority class
- would be nice to use all data for both training and testing, but we can't use training error directly due to overfitting
- Cross-validation to the rescue!


Basic idea

- we have a data set of n instances
- a model cannot be tested on its own training data due to overfitting, so need to divide into training and test set
- question we ask for each instance $i \in \{1, \dots, n\}$
 - "what would the model have predicted for the i :th instance, if it had not seen it during training?"
- we can answer this question by removing the i :th instance, training the model on the $n - 1$ remaining instances, and making a prediction for the i :th one (a test set of size 1)
- this is known as leave-one-out cross-validation

Leave-one-out cross-validation

- 1 **for** $i = 1$ to n **do**
- 2 train model on all but the i th instance
- 3 predict output for the i th instance with the model
- 4 **end for**
- 5 compute error between true and predicted outputs

Leave-one-out cross-validation



Analyzing leave-one-out cross-validation

- validating the models trained on $n - 1$ training instances gives us a good (almost unbiased) estimate of how the final model trained on all data will behave
- one each iteration, the model cannot overfit to the test instance since it is not part of training set
- leave-one-out allows us to use the whole data set for both training and testing simultaneously
- parameter selection: pick the model with lowest leave-one-out error (or highest AUC or correlation or...)
- final evaluation can also be done with leave-one-out

Implementing leave-one-out

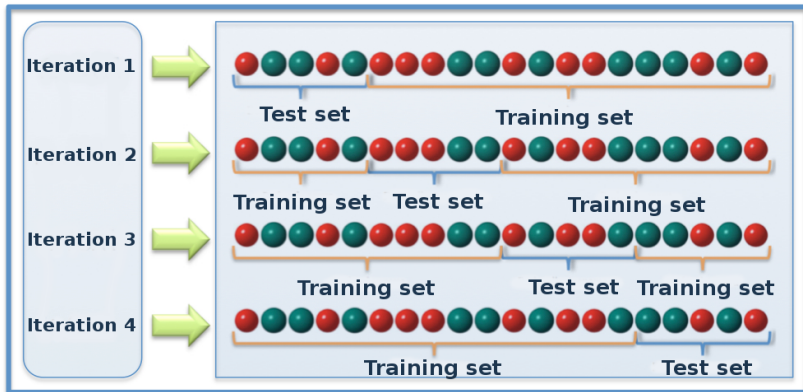
- for nearest neighbor extremely simple to implement
 - for each instance in training set, compute prediction from k nearest neighbors, excluding the instance itself
- generally, for other types of learning methods you need a loop where the machine learning method is trained n times
- can be computationally expensive
 - 5000 instances, 1 minute training time \rightarrow leave-one-out cross-validation takes 3.5 days
- Advanced topic: some machine learning methods allow computational shortcuts for computing leave-one-out much faster (k-nn, ridge regression, Naive Bayes...)

- for larger datasets leave-one-out can be impractical due to high computational costs
- **K-fold cross-validation**: same basic idea, but instead of leaving a single instance out at a time, leave several
- data is randomly divided into K (approximately) equally sized non-overlapping parts called **folds**
- on each round of cross-validation, use one fold for testing and remaining $K - 1$ for training
- usual choices for K : 10 or 5
- leave-one-out as extreme case, where $K = n$

K-fold cross-validation

- ① split the data randomly into K equally sized folds
- ② **for** $i = 1$ to K **do**
- ③ train model on data from all but the i th fold
- ④ predict outputs for the instances belonging to the i th fold with the model
- ⑤ **end for**
- ⑥ compute error between true and predicted outputs

K-fold cross-validation (K=4)



Analyzing K-fold cross-validation

- much faster than leave-one-out; 5000 instances, 1 minute training time → 10-fold CV takes 10 minutes
- small pessimistic bias, the model evaluated can be slightly less accurate than the final model trained on all data
- some variance depending on the random fold split
- rule of thumb: you can use leave-one-out on small data sets (e.g. 100 instances), 10-fold CV on medium sized (e.g. couple of thousands instances) and a single test set on large data sets

Some variants of K-fold cross-validation

- stratified K-fold cross-validation
 - especially imbalanced classification problems
 - stratification: split the data into folds so, that each fold roughly has the same fraction of members from each class as the full data set
 - guarantees that on all the iterations, the class distributions of the training and test set resemble that of the full data set
 - otherwise with very bad luck, it might happen that all instances from minority class end up in one fold
- M-times repeated K-fold cross-validation
 - some variance in the K-fold estimate due to the randomness in how the folds are split
 - can reduce the variance and thus obtain more reliable error estimates by repeating the whole procedure with different fold divisions M times (e.g. $M=100$) and taking the average
 - 5000 instances, 1 minute training time \rightarrow 100 times repeated 10-fold CV takes almost 17 hours (1000 minutes)

Combining model selection and final evaluation

- using cross-validation (or a single test set) we can do model selection
 - choose hyperparameter value / feature set / learning algorithm that gives lowest cross-validation error (or conversely, highest accuracy/AUC/correlation...)
- using cross-validation (or a single test set) we can evaluate the expected error rate of the final model
- what if we want to do both at the same time?

Combining model selection and final evaluation

- The error rate estimate of the final model biased (smaller than the true error rate) since out of several choices we pick the one having smallest estimated error
- example: 6 models, true expected error rates are $[0.40, 0.20, 0.10, 0.10, 0.10]$
- leave-one-out cross-validation gives us estimates: $[0.42, 0.18, 0.13, 0.10, 0.07]$
- we pick the last one, leading to a good choice of model, but the leave-one-out-estimate now biased (0.07 smaller than the true error rate 0.10)
- can be problematic especially on small datasets and when testing a large number of different models
- need additional independent test data, that was not used for model selection

Combining model selection and final evaluation

- first idea: instead of splitting data to training and test sets, split into three
- e.g. 50%, 25%, 25% split
- training set used for training the models to be evaluated
- we test all models against the validation set, pick one with lowest error
- final model is evaluated against the test set, which has not been used in any way until this time point
- good protocol if you have lots of data, can we combine this same idea with cross-validation?

Nested cross-validation

- nested, aka external cross-validation
- cross-validation within cross-validation
- inner loop for model selection, outer for evaluation
- on each round of the outer loop, data split to training and test set
 - compute a cross-validation estimate for each tested parameter on training set, choose best parameter
 - train on training set and predict on test set
- repeat for each fold

Nested cross-validation

- can be computationally very expensive (combinatorial explosion)
- leave-one-out within leave-one-out: need to train model $n * (n - 1) * (\text{\#tested params})$ times
- 5000 instances, 1 minute training time \rightarrow nested leave-one-out cross-validation takes about 48 years
- nested-leave-one-out feasible only on very small data sets, more often nested K -fold
- nesting especially important when you test a really large number of hyperparameters

- you should program cross-validation once by yourself to get the hang of it
- all major data analysis environments offer also this functionality
- see e.g. http://scikit-learn.org/stable/modules/cross_validation.html

Alternative criteria for model selection

- there are alternative approaches for model selection compared to using cross-validation
- information criteria, that balance fit to training data with model complexity
- e.g. Akaike's and the Bayesian information criterion (see book)
- general underlying principles, but the way the complexity can be evaluated is highly dependent on our assumptions about the type of models considered
- cross-validation more than competitive with these approaches, and generally applicable to any type of model used